

Project 2: Reinforcement Learning

Shreya S Ramanujam

AA228/CS238, Stanford University

SRAMANUJ@STANFORD.EDU

1. Algorithm Descriptions

1.1 Small Data Set and Large Data Set

For both the small as well as the large datasets, I used standard Q-learning, where the Q-table is updated directly from observed transitions given in the dataset (s, a, r, sp) and the final policy is greedy in Q-value with respect to all the actions at a particular state. However, for the larger datasets, many states were never visited in the given offline dataset. This was making direct Q-learning unstable, since using the Q-values of unvisited states from the table can cause unwanted error propagations in the Q value estimations.

To handle this, I modified the algorithm to interpolate Q-values for unseen states using their nearest visited counterparts. Specifically, I maintained a `visited_state` list to keep track of visited states, implying their Q-values are reliable. When updating or querying the Q-table for an unvisited state s , the algorithm locates the closest visited state based on state index (between 1 and 100 for small and 1 to 302020 for large) by doing a binary search on this `visited_state` list and uses its Q-values as a proxy in the Bellman update equation. This allows smoother propagation of information across sparsely sampled regions of the state space and a better overall policy. I used this interpolation both in the Bellman update equation (for both $Q(sp)$ and $Q(s)$ where needed) and also for the final greedy policy itself. The pseudocode is given in Algorithm 1.1.

Performance Characteristics: I ran the small dataset for **10** epochs, taking **4.44 seconds**. With small, even 10 epochs were sufficient to give a good policy. I ran the large dataset for **50** epochs, taking **46.53 seconds**. For large, I needed more epochs to arrive at a reasonable policy and estimate for Q, hence the larger number of epochs.

1.2 Medium Data Set

For the Medium Data Set, I tried the same algorithm detailed in 1.1, but I was never able to pass the autograder test. I also tried to include another clause in the Q-learning algorithm: if sp is a terminal state (≥ 0.45 in the original OpenAI Gymnasium documentation, which becomes State ≥ 457 after discretization to 500 states), then the target is just r , there is no additional $\gamma \max_a Q_{target}(a)$ term. However, this was not able to match the baseline either.

Finally, I settled on a degree 5 polynomial ridge regression to fit Q values: Poly(X, y). The predictor (X) values for the regression were the normalized [position, velocity] tuples and the targets were the Bellman update targets: $y = r + \gamma \max_a Q(sp, a)$. After fitting the polynomial, we use that to fill up our Q-table and greedily select actions for every state. The pseudocode is given in Algorithm 1.2.

Performance Characteristics: I ran this algorithm for **25** epochs, which took **27.98 seconds**. Any lesser number of epochs did not give a good policy (because the Q function was not well-estimated), which is why I chose to run for 25 epochs.

Algorithm 1 Q-Learning with Nearest-Visited-State Interpolation

```

1: Initialize  $Q(s, a) \leftarrow 0$  for all states  $s$  and actions  $a$ 
2: Initialize visited states list as empty
3: for each epoch do
4:   for each transition  $(s, a, r, sp)$  in shuffled data do
5:     if  $s$  not visited then
6:       Mark  $s$  as visited
7:        $Q_{update} \leftarrow Q(\text{nearest visited state}, .)$ 
8:     else
9:        $Q_{update} \leftarrow Q(s, .)$ 
10:    end if
11:    if  $sp$  not visited then
12:       $Q(sp, \cdot) \leftarrow Q(\text{nearest visited state}, \cdot)$ 
13:    else
14:       $Q_{target} \leftarrow Q(sp, \cdot)$ 
15:    end if
16:    Compute target:  $y \leftarrow r + \gamma \max_a Q_{target}(a)$ 
17:    Update:  $Q(s, a) \leftarrow Q_{update}(a) + \alpha(y - Q_{update}(a))$ 
18:  end for
19: end for
20: for each state  $s$  do
21:   if  $s$  visited then
22:      $\pi(s) \leftarrow \arg \max_a Q(s, a)$ 
23:   else
24:      $\pi(s) \leftarrow \arg \max_a Q(\text{nearest visited state}, a)$ 
25:   end if
26: end for
27: return policy  $\pi$ 

```

Algorithm 2 Polynomial Regression to find Q Values

```

1: Initialize  $Q(s, a) \leftarrow 0$  for all states  $s$  and actions  $a$ 
2: for each iteration do
3:   for each transition  $(s, a, r, s')$  in dataset do
4:     Compute target:  $y \leftarrow r + \gamma \max_{a'} Q(s', a')$ 
5:   end for
6:   Fit polynomial regression  $Q(s, a) \approx y$  using  $[s/n_{\text{states}}, a/n_{\text{actions}}]$  as features
7:   Predict  $Q(s, a)$  for all state-action pairs and update  $Q$ 
8: end for
9: Return policy:  $\pi(s) \leftarrow \arg \max_a Q(s, a)$ 

```

2. Code

```

1 import numpy as np
2 import pandas as pd
3 import torch.nn as nn
4 import argparse
5 import bisect
6 from sklearn.preprocessing import PolynomialFeatures
7 from sklearn.pipeline import make_pipeline
8 from sklearn.linear_model import Ridge
9 from time import time
10
11 ##### Helper functions, but I did not end up using them #####
12 def state_to_indices(position, velocity):
13     pos_idx = int((position + 1.2) / (0.6 + 1.2) * 499) # scale to [0,
14     ↵ 499]
15     vel_idx = int((velocity + 0.07) / (0.07 + 0.07) * 99) # scale to [0,
16     ↵ 99]
17     return pos_idx, vel_idx
18
19 def indices_to_state(pos_idx, vel_idx):
20     position = pos_idx / 499 * (0.6 + 1.2) - 1.2
21     velocity = vel_idx / 99 * (0.07 + 0.07) - 0.07
22     return position, velocity
23
24 def state_to_discretized_vals(state):
25     pos_idx = (state - 1) % 500
26     vel_idx = (state - 1) // 500
27     return (pos_idx, vel_idx)
28
29 def nearest_visited_state_mountaincar(visited_list, s):
30     pos_s , vel_s = state_to_discretized_vals(s + 1)
31     best_dist = float('inf')
32     nearest = None
33     for v in visited_list:
34         pos_v , vel_v = state_to_discretized_vals(v + 1)
35         dist = abs(pos_v - pos_s) + abs(vel_v - vel_s)
36         if dist < best_dist:
37             best_dist = dist
38             nearest = v
39     return nearest
40
41 ##### Unused Helper functions end #####
42
43 # function to find nearest visited neighbour for small and large

```

```

42 def nearest_visited_state(visited_list, state):
43     # binary search for nearest visited state
44     idx = bisect.bisect_left(visited_list, state)
45     if idx == 0:
46         nearest = visited_list[0]
47     elif idx == len(visited_list):
48         nearest = visited_list[-1]
49     else:
50         left = visited_list[idx - 1]
51         right = visited_list[idx]
52         nearest = left if abs(left - state) <= abs(right - state) else
53             ↵ right
54     return nearest
55
56 # code for medium dataset algorithm
57 def fitted_q_iteration_poly(data, n_states, n_actions, gamma=0.99,
58                             ↵ n_iters=50, degree=5):
59     model = make_pipeline(PolynomialFeatures(degree), Ridge(alpha=1.0))
60     q_values = np.zeros((n_states, n_actions))
61
62     for it in range(n_iters):
63         X, y = [], []
64         for _, row in data.iterrows():
65             s, a, r, sp = row['s'] - 1, row['a'] - 1, row['r'], row['sp'] -
66             ↵ 1
67             target = r + gamma * np.max(q_values[sp])
68             X.append([s / n_states, a / n_actions]) # to keep range same
69             y.append(target)
70
71         X, y = np.array(X), np.array(y)
72         model.fit(X, y)
73
74         sa_pairs = np.array([[s/n_states, a/n_actions] for s in
75             ↵ range(n_states) for a in range(n_actions)])
76         q_pred = model.predict(sa_pairs)
77         q_values = q_pred.reshape(n_states, n_actions)
78         print(f"Iteration {it+1}/{n_iters} done.")
79
80     return np.argmax(q_values, axis=1) + 1
81
82 # code for small and large dataset algorithm
83 def q_learning(task_name, data, n_states, n_actions, alpha=0.1, gamma=0.95,
84                 ↵ epochs=10):
85     # Initialize Q-table
86     q_table = np.zeros((n_states, n_actions)) # n_states and n_actions
87     # maintain a set of visited states

```

```

83     visited_states = np.zeros(n_states, dtype=bool)
84     visited_list = [] # keep sorted list of visited states
85     for epoch in range(epochs):
86         print(f"Epoch {epoch+1}/{epochs}")
87         for _, row in data.sample(frac=1, random_state=epoch).iterrows():
88             s = row['s'] - 1 # since states are 1-indexed in the data
89             a = row['a'] - 1 # since actions are 1-indexed in the data
90             r = row['r']
91             s_next = row['sp'] - 1 # since states are 1-indexed in the data
92             # for q_table, find nearest state that we have visited
93             # mark visited
94             if not visited_states[s]:
95                 bisect.insort(visited_list, s) # insert while keeping
96                 ↪ sorted
97                 visited_states[s] = True
98                 nearest = nearest_visited_state(visited_list, s)
99                 q_update = q_table[nearest]
100            else:
101                q_update = q_table[s]
102
102            if not visited_states[s_next]:
103                bisect.insort(visited_list, s_next) # insert while keeping
104                ↪ sorted
105                nearest = nearest_visited_state(visited_list, s_next)
106                q_target = q_table[nearest]
107            else:
108                q_target = q_table[s_next]
109
109            # if (task_name == 'medium' and state_to_discretized_vals(s +
110            ↪ 1)[0] >= 457): # terminal state
111            #     target = r
112            # else:
113            target = r + gamma * np.max(q_target)
114
114            q_table[s, a] = q_update[a] + alpha * (target - q_update[a])
115
116    policy = np.zeros(n_states, dtype=int)
117
118    # Derive policy using nearest visited state via binary search
119    for i in range(n_states):
120        if visited_states[i]:
121            policy[i] = np.argmax(q_table[i])
122        else:
123            nearest = nearest_visited_state(visited_list, i)
124            policy[i] = np.argmax(q_table[nearest])
125

```

```

126     return policy + 1 # convert back to 1-indexed
127
128
129 def main():
130     # get arguments from command line with defaults
131     parser = argparse.ArgumentParser()
132     parser.add_argument('--task_name', type=str, default='small',
133                         help='Task name: small, medium, large')
134     parser.add_argument('--alpha', type=float, default=0.1, help='Learning
135                         rate')
136     parser.add_argument('--gamma', type=float, default=None, help='Discount
137                         factor')
138     parser.add_argument('--num_epochs', type=int, default=None,
139                         help='Number of training epochs')
140     parser.add_argument('--degree', type=int, default=None, help='Degree of
141                         polynomial features for medium task')
142     args = parser.parse_args()
143
144     if(args.task_name == 'small'):
145         n_states = 100
146         n_actions = 4
147         input_data_path = './data/small.csv'
148         output_path = 'small.policy'
149         args.gamma = 0.95 if args.gamma is None else args.gamma
150         args.num_epochs = 10 if args.num_epochs is None else
151                         args.num_epochs
152
153     elif(args.task_name == 'medium'):
154         n_states = 50000
155         n_actions = 7
156         input_data_path = './data/medium.csv'
157         output_path = 'medium.policy'
158         args.gamma = 1.0 if args.gamma is None else args.gamma
159         args.num_epochs = 25 if args.num_epochs is None else
160                         args.num_epochs
161         args.degree = 5 if args.degree is None else args.degree
162
163     elif(args.task_name == 'large'):
164         n_states = 302020
165         n_actions = 9
166         input_data_path = './data/large.csv'
167         output_path = 'large.policy'
168         args.gamma = 0.95 if args.gamma is None else args.gamma
169         args.num_epochs = 50 if args.num_epochs is None else
170                         args.num_epochs

```

```

164     # read data
165     data = pd.read_csv(input_data_path)
166     # convert all data to ints
167     data = data.astype(int)
168     if args.task_name == 'medium':
169         print("Using specialized Q-learning for medium task")
170         start_time = time()
171         policy = fitted_q_iteration_poly(data, n_states, n_actions,
172             ↵ gamma=args.gamma, n_iters=args.num_epochs, degree=args.degree)
172         print(f"Medium task completed in {time() - start_time:.2f}
173             ↵ seconds")
173     else:
174         print("Using general Q-learning")
175         start_time = time()
176         policy = q_learning(args.task_name, data, n_states, n_actions,
177             ↵ alpha=args.alpha, gamma=args.gamma, epochs=args.num_epochs)
177         print(f"General Q-learning on {args.task_name} completed in {time()
178             ↵ - start_time:.2f} seconds")
178     # output to a file
179     with open(output_path, 'w') as f:
180         for _, action in enumerate(policy):
181             f.write(f"{action}\n")
182
183 if __name__ == "__main__":
184     main()
185

```
